

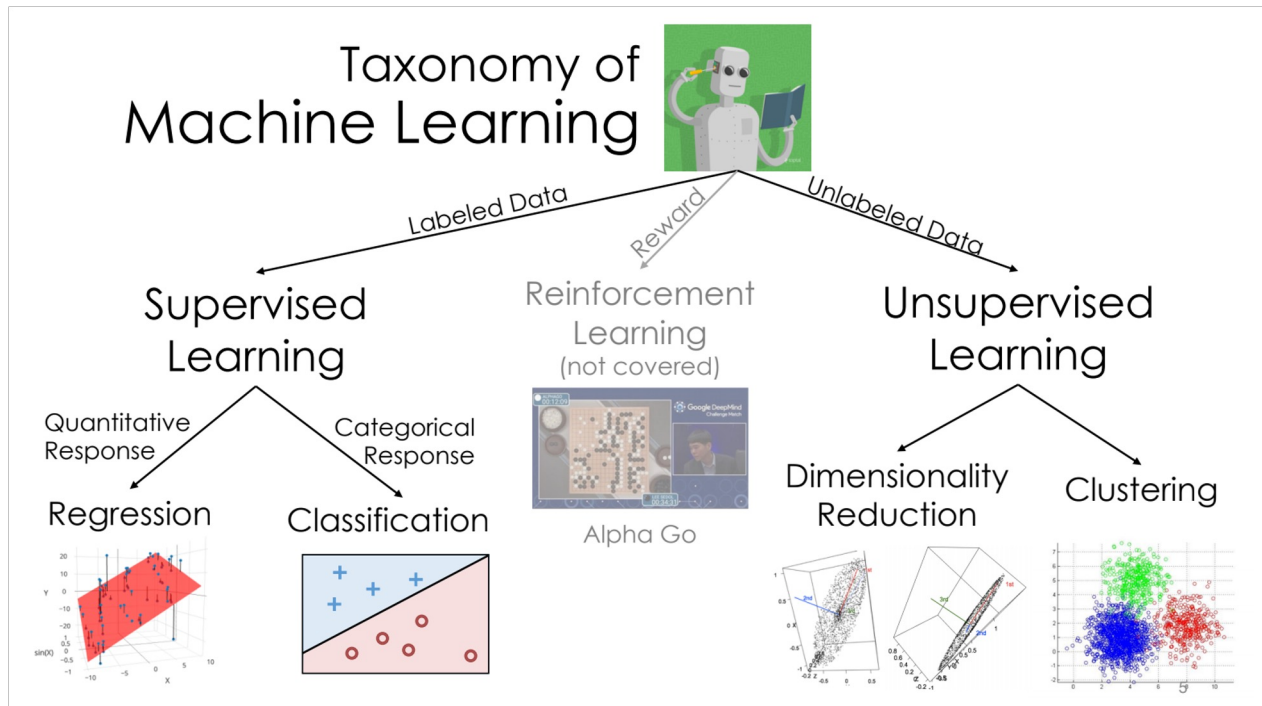
LECTURE 22

# Logistic Regression

Moving from regression to classification.

Regression and Classification are both forms of **supervised learning**.

**Logistic regression**, the topic of this lecture, is mostly used for **classification**, even though it has “regression” in the name.



# Regression vs Classification

---

## Regression vs Classification

Logistic regression model derivation

- logistic function (sigmoid)
- Parameter interpretation

Loss function

- Pitfalls of Squared Loss
- Cross Entropy

Maximum likelihood estimation

# So Far: Regression


In regression, we use unbounded numeric features to predict an *unbounded numeric output*.

GAME_ID	TEAM_NAME	MATCHUP	REB	FTM	TOV	GOAL_DIFF	WON
21700001	Boston Celtics	BOS @ CLE	46	19	12	-0.049	0
21700002	Golden State Warriors	GSW vs. HOU	41	19	17	0.053	0
21700003	Charlotte Hornets	CHA @ DET	47	23	17	-0.030	0
21700004	Indiana Pacers	IND vs. BKN	47	25	14	0.041	1
21700005	Orlando Magic	ORL vs. MIA	50	22	15	0.042	1

Input: numeric features

$$x^T \theta$$

Model: linear combination



Output: numeric prediction

Examples:

- Predict goal difference from turnover %
- Predict tip from total bill
- Predict mpg from hp

# Now: Classification

In **classification**, we use unbounded numeric features to predict a *categorical class*.

Examples:

- Predict *which team won* from turnover %
- Predict *day of week* from total bill
- Predict *model of car* from hp

	TEAM_NAME	MATCHUP	REB	FTM	TOV	GOAL_DIFF	WON
GAME_ID							
21700001	Boston Celtics	BOS @ CLE	46	19	12	-0.049	0
21700002	Golden State Warriors	GSW vs. HOU	41	19	17	0.053	0
21700003	Charlotte Hornets	CHA @ DET	47	23	17	-0.030	0
21700004	Indiana Pacers	IND vs. BKN	47	25	14	0.041	1
21700005	Orlando Magic	ORL vs. MIA	50	22	15	0.042	1

$$p = \sigma(x^\top \theta)$$

Win?  
If  $p > 0.5$ : predict a win  
Other: predict a loss



Input: numeric features

Model: linear combination transformed by non-linear **sigmoid**

**Decision rule**

Output: **class**

An aside: we will use logistic “regression” to perform a *classification* task. Here, “regression” refers to the type of model, not the task being performed.

# Kinds of Classification

We are interested in predicting some **categorical variable**, or **response**,  $y$ .

## Binary classification

- Two classes
- **Responses**  $y$  are either 0 or 1


win or lose

disease or no  
disease

spam or ham

## Multiclass classification

- Many classes
- Examples: Image labeling (Pishi, Thor, Hera), next word in a sentence, etc.



## Structured prediction tasks

- Multiple related classification predictions
- Examples: Translation, voice recognition, etc.

Our new goal: predict a **binary** output ( $y_{\text{hat}} = 0$  or  $y_{\text{hat}} = 1$ ) given inputted numeric features

## Regression ( $y \in \mathbb{R}$ )

### 1. Choose a model

Linear Regression

$$\hat{y} = f_{\theta}(x) = x^T \theta$$

### 2. Choose a loss function

Squared Loss or  
Absolute Loss

### 3. Fit the model

Regularization  
Sklearn/Gradient descent

### 4. Evaluate model performance

$R^2$ , Residuals, etc.

## Classification ( $y \in \{0, 1\}$ )

??

??

Regularization  
Sklearn/Gradient descent

??  
(next lecture)

# Logistic Regression model derivation

---

Regression vs Classification

## **Logistic regression model derivation**

- logistic function (sigmoid)
- Parameter interpretation

Loss function

- Pitfalls of Squared Loss
- Cross Entropy

Maximum likelihood estimation



## Example dataset

In this lecture, we will primarily use data from the 2017-18 NBA season.

**Goal:** Predict whether or not a team will win, given their FG\_PCT\_DIFF.

- This is the difference in field goal percentage between the two teams.
- Positive FG\_PCT\_DIFF: team made more shots than the opposing team.


TEAM_NAME	MATCHUP	WON	FG_PCT_DIFF
Boston Celtics	BOS @ CLE	0	-0.049
Golden State Warriors	GSW vs. HOU	0	0.053
Charlotte Hornets	CHA @ DET	0	-0.030
Indiana Pacers	IND vs. BKN	1	0.041
Orlando Magic	ORL vs. MIA	1	0.042

1s represent wins, 0s represent losses.

## Why not use Ordinary Least Squares?

We already have a model that can predict any quantitative response. Why not use it here?

- The output can be outside of the range  $[0, 1]$ . What does a predicted WON value of -2 mean?
- Very sensitive to outliers/ imbalanced data.
- Many other statistical reasons.




## Graph of averages

If we **bin the x-axis**, and take the **average** y-value for each bin, and tried to model that.

Doing so here yields a curve that resembles an S.

- Since our true y is either 0 or 1, this curve models the **probability that  $WON = 1$** , given  $FG\_PCT\_DIFF$ .
  - $WON = 1$  means “belong to class 1”.
- **Our goal is to model this red curve as best as possible.**



## Log-odds of probability is roughly linear

In the demo, we noticed that the **log-odds of the probability of belonging to class 1 was linear. This is the assumption that logistic regression is based on.**

$$\text{odds}(p) = \frac{p}{1-p} \quad \left| \quad \text{log-odds}(p) = \log \left( \frac{p}{1-p} \right) \right.$$

For now, let's let  $t$  denote our linear function (since log-odds is linear). Solving for  $p$ :

$$\begin{aligned} t &= \log \left( \frac{p}{1-p} \right) \\ e^t &= \frac{p}{1-p} \\ e^t - pe^t &= p \\ p &= \frac{e^t}{1+e^t} = \frac{1}{1+e^{-t}} \end{aligned}$$

With logistic regression, we are always referring to log base e ("ln").

## Log-odds of probability is roughly linear

In the demo, we noticed that the **log-odds of the probability of belonging to class 1 was linear. This is the assumption that logistic regression is based on.**

$$\text{odds}(p) = \frac{p}{1-p} \quad \left| \quad \text{log-odds}(p) = \log \left( \frac{p}{1-p} \right) \right.$$

For now, let's let  $t$  denote our linear function (since log-odds is linear). Solving for  $p$ :

$$\begin{aligned} t &= \log \left( \frac{p}{1-p} \right) \\ e^t &= \frac{p}{1-p} \\ e^t - pe^t &= p \\ p &= \frac{e^t}{1+e^t} = \frac{1}{1+e^{-t}} \end{aligned}$$

This is called the **logistic function**,  $\sigma(t)$ .

## Arriving at the logistic regression model

We know how to model linear functions quite well.

- We can substitute  $t = x^T \theta$ , since  $t$  was just a placeholder.

$p$  represents the probability of belonging to class 1.

- We are modeling  $P(Y = 1|x)$ .

$$p = \frac{1}{1 + e^{-t}} = \sigma(t)$$

Putting this all together:

$$P(Y = 1|x) = \frac{1}{1 + e^{-x^T \theta}} = \sigma(x^T \theta)$$

Looks just like the linear regression model, with a  $\sigma()$  wrapped around it.

We call logistic regression a **generalized linear model**, since it is a non-linear transformation of a linear model.

1. Choose a model

**Logistic Regression**

$$\hat{y} = f_{\theta}(x) = P(Y = 1|x) = \sigma(x^T \theta)$$

2. Choose a loss function

3. Fit the model

4. Evaluate model performance

## Linear vs. logistic regression

---

In a **linear regression** model, we predict a **quantitative** variable (i.e., some real number) as a linear function of features.

- Our output, or **response**,  $y$ , could be any real number.

$$\hat{y} = f_{\theta}(x) = x^T \theta$$

In a **logistic regression** model, our goal is to predict a binary **categorical** variable (class 0 or class 1) as a linear function of features, passed through the logistic function.

- Our **response** is the probability that our observation belongs to class 1.
- Haven't yet done classification!

$$\hat{y} = f_{\theta}(x) = P(Y = 1|x) = \sigma(x^T \theta)$$



# Logistic Function

---

Regression vs Classification

Logistic regression model derivation

- **logistic function (sigmoid)**
- Parameter interpretation

Loss function

- Pitfalls of Squared Loss
- Cross Entropy


Maximum likelihood estimation

## Properties of the logistic function

The logistic function is a type of **sigmoid**, a class of functions that share certain properties.

$$\sigma(t) = \frac{1}{1 + e^{-t}} \quad -\infty < t < \infty$$

- Its output is bounded between 0 and 1, no matter how large  $t$  is.
  - Fixes an issue with using linear regression to predict probabilities.
- We can interpret it as mapping real numbers to probabilities.



# Properties of the logistic function

---

## Definition

$$\sigma(t) = \frac{1}{1 + e^{-t}} = \frac{e^t}{1 + e^t}$$

## Range

$$0 < \sigma(t) < 1$$

## Inverse

$$t = \sigma^{-1}(p) = \log\left(\frac{p}{1-p}\right)$$

## Reflection and Symmetry

$$1 - \sigma(t) = \frac{e^{-t}}{1 + e^{-t}} = \sigma(-t)$$

## Derivative


$$\frac{d}{dt}\sigma(t) = \sigma(t)(1 - \sigma(t)) = \sigma(t)\sigma(-t)$$

## Shape of the logistic function

Consider the plot of  $\sigma(\theta_1 x)$ , for several different values of  $\theta_1$ .

- If  $\theta_1$  is positive, the curve increases to the right.
- The further  $\theta_1$  is from 0, the steeper the curve.

In the notebook, we explore more sophisticated logistic curves.



# Parameter interpretation

---

Regression vs Classification

Logistic regression model derivation

- logistic function (sigmoid)
- **Parameter interpretation**

Loss function

- Pitfalls of Squared Loss
- Cross Entropy

Maximum likelihood estimation

## Parameter interpretation

Recall, we arrived at the model by assuming that the **log-odds of the probability of belonging to class 1 was linear**.

$$P(Y = 1|x) = \sigma(x^T \theta) \quad \leftarrow \quad \log \left( \frac{P(Y = 1|x)}{P(Y = 0|x)} \right) = x^T \theta \quad \leftarrow \quad \frac{P(Y = 1|x)}{P(Y = 0|x)} = e^{x^T \theta}$$

This is the same as  $\frac{p}{1-p}$  because

$$P(Y = 1|x) + P(Y = 0|x) = 1$$

(Remember, we are dealing with binary classification – we are predicting 1 or 0.)

## Parameter interpretation

Let's suppose our linear component has just a single feature, along with an intercept term.

$$\frac{P(Y = 1|x)}{P(Y = 0|x)} = e^{\theta_0 + \theta_1 x}$$

What happens if you increase  $x$  by one unit?

- Odds is multiplied by  $e^{\theta_1}$ .
- If  $\theta_1 > 0$ , the odds increase.
- If  $\theta_1 < 0$ , the odds decrease.

What happens if  $x^T \theta = \theta_0 + \theta_1 x = 0$ ?

- This means class 1 and class 0 are equally likely.

- $$e^0 = 1 \implies \frac{P(Y = 1|x)}{P(Y = 0|x)} = 1 \implies P(Y = 1|x) = P(Y = 0|x)$$

The odds ratio can be interpreted as the “number of successes for each failure.”

# Loss Function

---

Regression vs Classification

Logistic regression model derivation

- logistic function (sigmoid)
- Parameter interpretation

## Loss Function

- Pitfalls of Squared Loss
- Cross Entropy

Maximum likelihood estimation



# Logistic Regression with squared loss?

1. Choose a model

**Logistic Regression**

$$\hat{y} = f_{\theta}(x) = P(Y = 1|x) = \sigma(x^T \theta)$$

2. Choose a loss function

$$R(\theta) = \frac{1}{n} \sum_{i=1}^n (y_i - \sigma(\mathbb{X}_i^T \theta))^2 \quad ?$$

3. Fit the model


4. Evaluate model performance

## Pitfalls of squared loss with logistic regression


On the left, we have a toy dataset (i.e. we've plotted the original data,  $y$  vs.  $x$ ).

On the right, we have a plot of the mean squared error of this dataset when fitting a single-feature logistic regression model, for different values of  $\theta$  (i.e. the loss surface).

What is the issue?



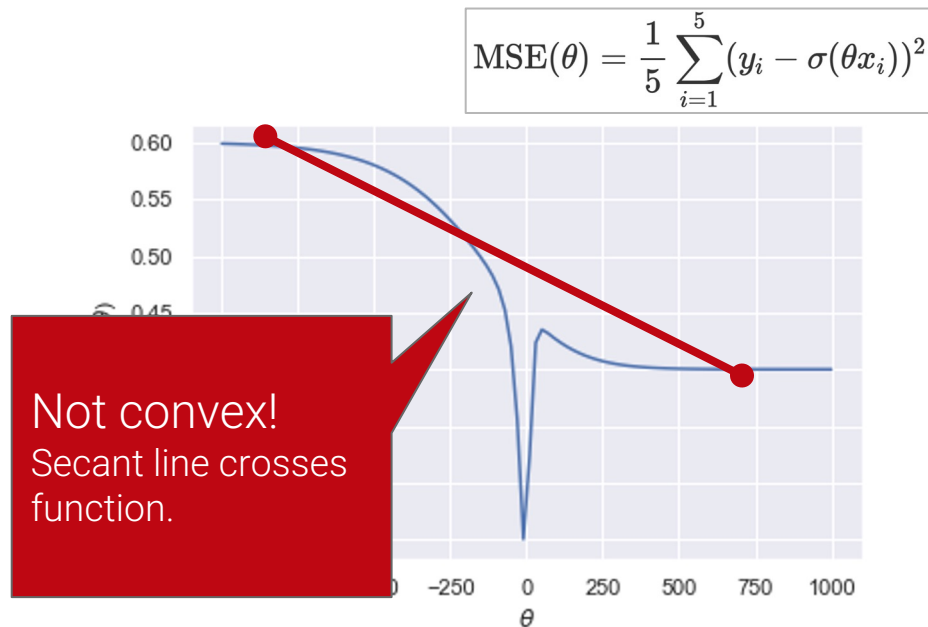

$$\text{MSE}(\theta) = \frac{1}{5} \sum_{i=1}^5 (y_i - \sigma(\theta x_i))^2$$



## Pitfalls of squared loss with logistic regression

On the **left**, we have a toy dataset (i.e. we've plotted the **original data**,  $y$  vs.  $x$ ).

On the **right**, we have a plot of the **mean squared error** of this dataset when fitting a single-feature logistic regression model, for different values of  $\theta$  (i.e. the **loss surface**).



## Pitfalls of squared loss with logistic regression

For this particular loss surface, different initial guesses for  $\theta$  yield different “optimal values”, as per `scipy.optimize.minimize`:


```
1 minimize(mse_loss_single_arg_toy, x0 = 0) ["x"][0]
```

-4.801981341432673

```
1 minimize(mse_loss_single_arg_toy, x0 = 500) ["x"][0]
```

500.0

This loss surface is not convex. We'd like it to be.




## Pitfalls of squared loss with logistic regression

The loss surface of MSE for a logistic regression model with a single slope plus an intercept often looks something like this.

If your initial guess for  $\hat{\theta}$  is way out in the flat yellow region, routine can get stuck. The model gets stuck in local minima

If the gradient is 0, your update rule will stop changing.

$$\theta^{(t+1)} = \theta^{(t)} - \alpha \nabla_{\theta} R(\theta, \mathbb{X}, \mathbb{Y})$$




## Pitfalls of squared loss with logistic regression

Another issue: since  $y_i$  is either 0 or 1, and  $\hat{y}_i$  is between 0 and 1,  $(y_i - \hat{y}_i)^2$  is also bounded between 0 and 1.

- Even if our probability is nowhere close, the loss isn't that large in magnitude.
  - If we say the probability is  $10^{-6}$ , but it happens anyway, error should be large.
- We want to penalize wrong answers significantly.

Suppose the observation we're trying to predict is actually in **class 1**.

On the right, we have a plot of  $(1 - \hat{y})^2$  vs  $\hat{y}$ . This is the squared loss for a single prediction.



## Summary of issues with squared loss and logistic regression

---

While it can work, squared loss is not the best choice of loss function for logistic regression.

- Average squared loss is not convex.
  - Numerical methods will struggle to find a solution.
- Wrong predictions aren't penalized significantly enough.
  - Squared loss (and hence, average squared loss) are bounded between 0 and 1.

Fortunately, there's a solution.

## Loss in Classification

---

Let  $y$  be a binary label  $\{0, 1\}$ , and  $p$  be the model's predicted probability of the label being 1.

In a classification task, how do we want our loss function to behave?

- When the true  $y$  is 1, we should incur low loss when the model predicts large  $p$ .
- When the true  $y$  is 0, we should incur high loss when the model predicts large  $p$ .

In other words, the behavior we need from our loss function depends on the value of the true class,  $y$ .




# Cross-Entropy Loss

Let  $y$  be a binary label  $\{0, 1\}$ , and  $p$  be the probability of the label being 1.


The **cross-entropy loss** is defined as:

$$\text{CE loss} = \begin{cases} -\log(p), & \text{if } y = 1 \\ -\log(1 - p), & \text{if } y = 0 \end{cases}$$



For  $y = 1$ ,

- $p \rightarrow 0$ :  $\infty$  loss
- $p \rightarrow 1$ : zero loss



For  $y = 0$ ,

- $p \rightarrow 0$ : zero loss
- $p \rightarrow 1$ :  $\infty$  loss

# Cross-Entropy Loss: Two Loss Functions In One!

The piecewise loss function we introduced just then is difficult to optimize – we don't want to check “which” loss to use at each step of optimizing theta.

Cross-entropy loss can be equivalently expressed as:



$$\text{CE loss} = \begin{cases} -\log(p), & \text{if } y = 1 \\ -\log(1 - p), & \text{if } y = 0 \end{cases}$$

$$\rightarrow - (y \log(p) + (1 - y) \log(1 - p))$$

makes loss positive

for  $y = 1$ , only this term stays

for  $y = 0$ , only this term stays



## Empirical Risk: Average Cross-Entropy Loss

For a single datapoint, the cross-entropy curve is convex. It has a global minimum.

$$-(y \log(p) + (1 - y) \log(1 - p))$$

What about average cross-entropy loss, i.e., empirical risk?

For logistic regression, the empirical risk over a sample of size  $n$  is:


$$\begin{aligned} R(\theta) &= -\frac{1}{n} \sum_{i=1}^n (y_i \log(p_i) + (1 - y_i) \log(1 - p_i)) \\ &= -\frac{1}{n} \sum_{i=1}^n (y_i \log(\sigma(X_i^T \theta)) - (1 - y_i) \log(1 - \sigma(X_i^T \theta))) \quad [\text{Recall our model is } \hat{y}_i = p_i = \sigma(X_i^T \theta)] \end{aligned}$$

The optimization problem is therefore to find the estimate  $\hat{\theta}$  that minimizes  $R(\theta)$ :


$$\hat{\theta} = \underset{\theta}{\operatorname{argmin}} -\frac{1}{n} \sum_{i=1}^n (y_i \log(\sigma(X_i^T \theta)) - (1 - y_i) \log(1 - \sigma(X_i^T \theta)))$$

# Convexity Proof By Picture


$$\hat{\theta} = \underset{\theta}{\operatorname{argmin}} \quad - \frac{1}{n} \sum_{i=1}^n \left( y_i \log \left( \sigma \left( X_i^T \theta \right) \right) - (1 - y_i) \log \left( 1 - \sigma \left( X_i^T \theta \right) \right) \right)$$



	x	y
0	-4.0	0
1	-2.0	0
2	-0.5	1
3	1.0	0
4	3.0	1
5	5.0	1



A straight line crosses the curve  
Non-convex



Convex!

1. Choose a model

**Logistic Regression**

$$\hat{y} = f_{\theta}(x) = P(Y = 1|x) = \sigma(x^T \theta)$$

2. Choose a loss function

$$\cancel{R(\theta) = \frac{1}{n} \sum_{i=1}^n (y_i - \sigma(\mathbb{X}_i^T \theta))^2} \quad \text{loss} = -y \log(\hat{y}) - (1 - y) \log(1 - \hat{y})$$

3. Fit the model

$$R(\theta) = -\frac{1}{n} \sum_{i=1}^n (y_i \log(\sigma(\mathbb{X}_i^T \theta)) + (1 - y_i) \log(1 - \sigma(\mathbb{X}_i^T \theta)))$$

4. Evaluate model performance

1. Choose a model

**Logistic Regression**

$$\hat{y} = f_{\theta}(x) = P(Y = 1|x) = \sigma(x^T \theta)$$

2. Choose a loss function

$$R(\theta) = -\frac{1}{n} \sum_{i=1}^n (y_i \log(\sigma(\mathbb{X}_i^T \theta)) + (1 - y_i) \log(1 - \sigma(\mathbb{X}_i^T \theta)))$$

3. Fit the model

For logistic regression, we can use squared loss if we want to!

- Using squared loss and using cross-entropy loss will usually result in different  $\hat{\theta}$ 
  - **Different optimization problems, different solutions.**
- **Cross-entropy loss is strictly better than squared loss for logistic regression.**
  - Convex, so easier to minimize using numerical techniques.
  - Better suited for modeling probabilities.

4. Evaluate model performance

# Maximum Likelihood Estimation

---

Regression vs Classification

Logistic regression model derivation

- logistic function (sigmoid)
- Parameter interpretation

Loss function

- Pitfalls of Squared Loss
- Cross Entropy

**Maximum Likelihood Estimation**

## PMF of the Bernoulli distribution

---

If  $Y$  is the result of one toss of a coin that lands heads with chance  $p$ ,

$$\begin{aligned} P(Y = y) &= \begin{cases} p & \text{if } y = 1 \\ 1 - p & \text{if } y = 0 \end{cases} \\ &= p^y (1 - p)^{1-y} \quad \triangleleft \end{aligned}$$

Then, if  $Y_1$  and  $Y_2$  are the results of tosses of two coins:

$$P(Y_1 = y_1, Y_2 = y_2) = (p_1^{y_1} (1 - p_1)^{1-y_1}) (p_2^{y_2} (1 - p_2)^{1-y_2})$$

---



## Estimating the two probabilities

---

- Suppose we want to estimate the values of  $p_1$  and  $p_2$ .
- We know what the likelihood is.

$$P(Y_1 = y_1, Y_2 = y_2) = (p_1^{y_1} (1 - p_1)^{1-y_1}) (p_2^{y_2} (1 - p_2)^{1-y_2})$$

---

- Our goal is to find the  $p_1$  and  $p_2$  that **maximize** the above function, over all  $p_1$  and  $p_2$ .
  - Maximize, because we are looking for the  $p_1$  and  $p_2$  that are “most likely” to have generated the data that we observed.
- As before, this involves differentiating, setting equal to 0, and solving.

## Log likelihoods

- Maximizing  $P(Y_1 = y_1, Y_2 = y_2) = (p_1^{y_1}(1 - p_1)^{1-y_1})(p_2^{y_2}(1 - p_2)^{1-y_2})$  is annoying.
  - Products  $\rightarrow$  chain rule.
- $\log(x)$  is a **strictly increasing** function.
  - If  $a > b$ , then  $\log(a) > \log(b)$ .
- This means, the values of  $p_1$  and  $p_2$  that maximize  $P(Y_1 = y_1, Y_2 = y_2)$  are the same values that maximize

$$\begin{aligned} & \log \left( (p_1^{y_1}(1 - p_1)^{1-y_1})(p_2^{y_2}(1 - p_2)^{1-y_2}) \right) \\ &= y_1 \log(p_1) + (1 - y_1) \log(1 - p_1) + y_2 \log(p_2) + (1 - y_2) \log(1 - p_2) \\ &= \sum_{i=1}^2 (y_i \log(p_i) + (1 - y_i) \log(1 - p_i)) \end{aligned}$$

Starting to look familiar!

## Estimating $n$ probabilities

- For  $i = 1, 2, \dots, n$ , let  $Y_i$  be Bernoulli( $p_i$ ).
  - Each  $Y_i$  is independent of each other.
- To estimate  $p_1, p_2, \dots, p_n$ :

Find  $\hat{p}_1, \hat{p}_2, \dots, \hat{p}_n$  that maximize  $\sum_{i=1}^n (y_i \log(p_i) + (1 - y_i) \log(1 - p_i))$

**Equivalently:**

Find  $\hat{p}_1, \hat{p}_2, \dots, \hat{p}_n$  that *minimize*  $-\frac{1}{n} \sum_{i=1}^n (y_i \log(p_i) + (1 - y_i) \log(1 - p_i))$

We choose this equivalent form because we are more used to minimizing loss.

## Cross-entropy loss

Find  $\hat{p}_1, \hat{p}_2, \dots, \hat{p}_n$  that *minimize*  $-\frac{1}{n} \sum_{i=1}^n (y_i \log(p_i) + (1 - y_i) \log(1 - p_i))$

### What does this have to do with logistic regression?

- We have  $n$  observations  $(y_1, y_2, \dots, y_n)$ . Each is either 1 or 0.
  - Assume that each is independent of one another.
- Can think of observation  $y_i$  as the result of a coin toss with probability  $p_i$ .
  - The output of our logistic regression model is our estimate for the probability that  $y_i = 1$ .
- Thus, we can use the above average loss, with  $p_i = \sigma(\mathbb{X}_i^T \theta)$ .
- This gives us the exact same expression for cross-entropy loss that we saw before!

## Maximum likelihood estimation

---

Minimizing cross-entropy loss is equivalent to maximizing the likelihood of the data.

- We are choosing the model parameters that are “most likely”, given this data.
- Another perspective of fitting our model to the data.

$$R(\theta) = -\frac{1}{n} \sum_{i=1}^n (y_i \log(\sigma(\mathbb{X}_i^T \theta)) + (1 - y_i) \log(1 - \sigma(\mathbb{X}_i^T \theta)))$$

This technique is called **maximum likelihood estimation (MLE)**.

- It turns out, many of the model + loss combinations we’ve seen in this class can be motivated using MLE.
  - OLS, Ridge Regression.
- You will study this further in probability and ML classes. But now you know it exists.

$$R(\theta) = -\frac{1}{n} \sum_{i=1}^n (y_i \log(\sigma(\mathbb{X}_i^T \theta)) + (1 - y_i) \log(1 - \sigma(\mathbb{X}_i^T \theta)))$$

- Approach 1: Gradient Descent (take larger – more certain – steps opposite the gradient)
- Approach 2: Stochastic Gradient Descent (SGD) (take many small steps opposite the gradient)
- Approach 3: Newton's Method (use second derivatives to better follow curvature)

1. Choose a model

**Logistic Regression**

$$\hat{y} = f_{\theta}(x) = P(Y = 1|x) = \sigma(x^T \theta)$$

2. Choose a loss function

$$R(\theta) = -\frac{1}{n} \sum_{i=1}^n (y_i \log(\sigma(\mathbb{X}_i^T \theta)) + (1 - y_i) \log(1 - \sigma(\mathbb{X}_i^T \theta)))$$

3. Fit the model



4. Evaluate model performance

Next time!